

Import a JSON file from the command line and apply actions
with the data present in the JSON file

Aim:

To import a JSON file from the command line and apply the following actions with the data present in the JSON file where, projection, aggregation, remove, count, limit, skip and sort.

Procedure:**Hive Download and installation:**

1. Starting Hadoop Services

Open PowerShell as administrator and go to Hadoop sbin directory and start hadoop services using the following commands: `Start-all.cmd`

```
C:\Windows\System32>start-all.cmd
This script is Deprecated. Instead use start-dfs.cmd and start-yarn.cmd
starting yarn daemons
```

2. Create a .json file with the below content:

```
{ "id": 1, "name": "John Doe", "age": 30, "salary": 50000 }
{ "id": 2, "name": "Jane Smith", "age": 25, "salary": 60000 }
{ "id": 3, "name": "Alice Johnson", "age": 28, "salary": 55000 }
{ "id": 4, "name": "Bob Brown", "age": 35, "salary": 70000 }
{ "id": 5, "name": "Charlie Davis", "age": 40, "salary": 80000 }
{ "id": 6, "name": "Eve White", "age": 22, "salary": 48000 }
{ "id": 7, "name": "Frank Black", "age": 32, "salary": 65000 }
{ "id": 8, "name": "Grace Green", "age": 27, "salary": 52000 }
{ "id": 9, "name": "Henry Gold", "age": 29, "salary": 59000 } { "id":
10, "name": "Isabel Blue", "age": 33, "salary": 73000 } Derby
```

Network Server:

Run the following command to open Derby:

```
StartNetworkServer -h 0.0.0.0
```

```
C:\Windows\System32>startNetworkServer -h 0.0.0.0
Fri Sep 13 19:17:50 IST 2024 : Security manager installed using the Basic server security policy.
Fri Sep 13 19:17:50 IST 2024 : Apache Derby Network Server - 10.14.2.0 - (1828579) started and ready to accept connections on port 1527
-
```

Go to first PowerShell window and check whether NetworkServerControl is running.

```
C:\Windows\System32>jps
2976 DataNode
9392 NodeManager
11640 NetworkServerControl
11116 NameNode
23452 Jps
2668 ResourceManager
C:\Windows\System32>
```

3. Starting Apache Hive:

Go to Apache Hive's bin location with cd command and run the following command:

```
hive --service schematool -dbType derby --initSchema
```

8. Open Hive shell by typing:

```
hive
```

```
C:\hive\bin>hive
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/C:/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.30.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/C:/hive/lib/log4j-slf4j-impl-2.17.1.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
2024-09-13 19:33:54,247 INFO conf.HiveConf: Found configuration file null
2024-09-13 19:33:55,244 WARN common.LogUtils: hive-site.xml not found on CLASSPATH
```

Create a Database:

Start by creating a database. Open the Hive CLI and follow the steps below:

1. Use the **CREATE DATABASE** statement to create a new database:

```
CREATE DATABASE IF NOT EXISTS emp_json;
```

2. Verify the database is present:

```
SHOW DATABASES;
```

```
hive> SHOW DATABASES;
2024-09-13 19:37:28,134 INFO conf.HiveConf: Using the default value passed in for log id: 041a6a49-576e-4335-bb38-a8565534cf6b
2024-09-13 19:37:28,134 INFO session.SessionState: Updating thread name to 041a6a49-576e-4335-bb38-a8565534cf6b main
2024-09-13 19:37:28,136 INFO ql.Driver: Compiling command(queryId=monid_20240913193728_7db73a2b-da6c-46a5-bb28-4618c09fa0d8): SHOW DATABASES
2024-09-13 19:37:28,186 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-13 19:37:28,206 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-13 19:37:28,270 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:database_name, type:string, comment:from deserializer)], properties:null)
2024-09-13 19:37:28,357 INFO exec.ListSinkOperator: Initializing operator LIST_SINK[0]
2024-09-13 19:37:28,366 INFO ql.Driver: Completed compiling command(queryId=monid_20240913193728_7db73a2b-da6c-46a5-bb28-4618c09fa0d8); Time taken: 0.23 seconds
2024-09-13 19:37:28,367 INFO reexec.ReExecDriver: Execution #1 of query
2024-09-13 19:37:28,368 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-13 19:37:28,368 INFO ql.Driver: Executing command(queryId=monid_20240913193728_7db73a2b-da6c-46a5-bb28-4618c09fa0d8): SHOW DATABASES
2024-09-13 19:37:28,380 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-09-13 19:37:28,381 INFO metastore.HiveMetaStore: 0: get_databases: @hive#
2024-09-13 19:37:28,382 INFO HiveMetaStore.audit: ugi=monid ip=unknown-ip-addr cmd=get_databases: @hive#
2024-09-13 19:37:28,384 INFO exec.DDLTask: results : 2
```

Create a Table in Hive:

```
CREATE TABLE employees_table (
```

```
    id INT,
```

```
    name STRING,
```

```
    age INT, salary
```

```
    DOUBLE
```

)

ROW FORMAT SERDE 'org.apache.hive.hcatalog.data.JsonSerDe'

STORED AS TEXTFILE

LOCATION '/ser/hive/warehouse/emp_json/';

```
2024-09-13 18:41:24,255 INFO session.SessionState: Resetting thread name to main
hive> CREATE TABLE employees_table (
  > id INT,
  > name STRING,
  > age INT,
  > salary DOUBLE
  > )
  > ROW FORMAT SERDE 'org.apache.hive.hcatalog.data.JsonSerDe'
  > STORED AS TEXTFILE
  > LOCATION '/user/hive/warehouse/emp_json/';
2024-09-13 18:42:28,011 INFO conf.HiveConf: Using the default value passed in for log id: 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1
2024-09-13 18:42:28,011 INFO session.SessionState: Updating thread name to 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1 main
2024-09-13 18:42:28,013 INFO ql.Driver: Compiling command(queryId=monid_20240913184228_e23cbc73-eea9-417c-b9f0-f206513f3b35): CREATE TABLE employees_table (
id INT,
name STRING,
age INT,
salary DOUBLE
)
ROW FORMAT SERDE 'org.apache.hive.hcatalog.data.JsonSerDe'
STORED AS TEXTFILE
LOCATION '/user/hive/warehouse/emp_json/'
```

Add Data to the TABLE:

Run the **LOAD DATA INPATH** command:

```
LOAD DATA INPATH '/user/hive/warehouse/emp_json/employee.json' INTO TABLE employees_table;
```

```
hive> LOAD DATA INPATH '/user/hive/warehouse/employee.json' INTO TABLE employees_table;
2024-09-13 18:43:03,505 INFO conf.HiveConf: Using the default value passed in for log id: 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1
2024-09-13 18:43:03,505 INFO session.SessionState: Updating thread name to 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1 main
2024-09-13 18:43:03,508 INFO ql.Driver: Compiling command(queryId=monid_20240913184303_e143d92f-b894-4b56-855c-d687e518f311): LOAD DATA INPATH '/user/hive/warehouse/employee.json' INTO TABLE employees_table
2024-09-13 18:43:03,519 INFO metastore.HiveMetaStoreClient: Metastore configuration metastore.filter.hook changed from org.apache.hadoop.hive.metastore.DefaultMetaStoreFilterHookImpl to org.apache.hadoop.hive.ql.security.authorization.plugin.AuthorizationMetaStoreFilterHook
2024-09-13 18:43:03,520 INFO metastore.HiveMetaStore: 0: Cleaning up thread local RawStore...
2024-09-13 18:43:03,523 INFO HiveMetaStore.audit: ugi=monid ip=unknown-ip-addr cmd=Cleaning up thread local RawStore
```

List Hive Tables and Data:

To show all tables in a selected database, use the following statement:

SHOW TABLES;

```
2024-09-13 18:43:19,251 INFO session.SessionState: Resetting thread name to main
hive> SHOW TABLES;
2024-09-13 18:43:19,246 INFO conf.HiveConf: Using the default value passed in for log id: 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1
2024-09-13 18:43:19,246 INFO session.SessionState: Updating thread name to 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1 main
2024-09-13 18:43:19,249 INFO ql.Driver: Compiling command(queryId=monid_20240913184319_786f1cd8-47e7-4cd5-bc39-48276b96a9fc): SHOW TABLES
2024-09-13 18:43:19,260 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-13 18:43:19,262 INFO metastore.HiveMetaStore: 0: get_database: @hive#default
2024-09-13 18:43:19,262 INFO HiveMetaStore.audit: ugi=monid ip=unknown-ip-addr cmd=get_database: @hive#default
2024-09-13 18:43:19,264 INFO ql.Driver: Semantic Analysis Completed (retrial = false)
2024-09-13 18:43:19,264 INFO ql.Driver: Returning Hive schema: Schema(fieldSchemas:[FieldSchema(name:tab_name, type:string, comment:from deserializer)], properties:null)
2024-09-13 18:43:19,265 INFO exec.ListSinkOperator: Initializing operator LIST_SINK[0]
```

To show table column names and data types, run:

DESC employees_table;

```
ng, comment:from deserializer), FieldSchema(name:data_type, type:string, comment:from deserializer), FieldSchema(name:comment, type:string, comment:from deserializer)], properties:null)
2024-09-13 18:43:43,120 INFO exec.ListSinkOperator: Initializing operator LIST_SINK[0]
2024-09-13 18:43:43,121 INFO ql.Driver: Completed compiling command(queryId=monid_20240913184343_d830a95f-90fb-4f82-9f75-361591c07d0a); Time taken: 0.028 seconds
2024-09-13 18:43:43,121 INFO reexec.ReExecDriver: Execution #1 of query
2024-09-13 18:43:43,122 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-13 18:43:43,122 INFO ql.Driver: Executing command(queryId=monid_20240913184343_d830a95f-90fb-4f82-9f75-361591c07d0a): DESC employees_table
2024-09-13 18:43:43,125 INFO ql.Driver: Starting task [Stage-0:DDL] in serial mode
2024-09-13 18:43:43,126 INFO metastore.HiveMetaStore: 0: get_table : tbl=hive.default.employees_table
2024-09-13 18:43:43,126 INFO HiveMetaStore.audit: ugi=monid ip=unknown-ip-addr cmd=get_table : tbl=hive.default.employees_table
2024-09-13 18:43:43,144 INFO ql.Driver: Completed executing command(queryId=monid_20240913184343_d830a95f-90fb-4f82-9f75-361591c07d0a); Time taken: 0.022 seconds
OK
2024-09-13 18:43:43,145 INFO ql.Driver: OK
2024-09-13 18:43:43,147 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-13 18:43:43,150 INFO mapred.FileInputFormat: Total input files to process : 1
2024-09-13 18:43:43,173 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_0:4,
id                int                from deserializer
name              string            from deserializer
age              int              from deserializer
salary           double           from deserializer
Time taken: 0.054 seconds, Fetched: 4 row(s)
2024-09-13 18:43:43,179 INFO CliDriver: Time taken: 0.054 seconds, Fetched: 4 row(s)
```

To display table data, use a **SELECT** statement. For example, to select everything in a table, run:

```
SELECT * FROM employees_table;
```

```
64b): SELECT * FROM employees_table
2024-09-13 18:43:58,170 INFO ql.Driver: Completed executing command(queryId=monid_20240913184356_d209bc3b-0c89-4c16-af8c-9e7116da564b); Time taken: 0.001 seconds
OK
2024-09-13 18:43:58,171 INFO ql.Driver: OK
2024-09-13 18:43:58,172 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-13 18:43:58,180 INFO mapred.FileInputFormat: Total input files to process : 1
2024-09-13 18:43:58,200 INFO exec.TableScanOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_TS_0:10,
2024-09-13 18:43:58,200 INFO exec.SelectOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_SEL_1:10,
2024-09-13 18:43:58,200 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_3:10,
1   John Doe      30      50000.0
2   Jane Smith    25      60000.0
3   Alice Johnson 28      55000.0
4   Bob Brown     35      70000.0
5   Charlie Davis 40      80000.0
6   Eve White     22      48000.0
7   Frank Black   32      65000.0
8   Grace Green   27      52000.0
9   Henry Gold    29      59000.0
10  Isabel Blue   33      73000.0
Time taken: 1.395 seconds, Fetched: 10 row(s)
2024-09-13 18:43:58,219 INFO CliDriver: Time taken: 1.395 seconds, Fetched: 10 row(s)
```

Perform Various Operations on the Data in the table:

WHERE:

```
FROM SELECT id, name, age, salary employees_table
WHERE salary > 60000;
```

```
hive> SELECT id, name, age, salary FROM employees_table WHERE salary > 60000;
2024-09-13 18:54:02,186 INFO conf.HiveConf: Using the default value passed in for log id: 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1
2024-09-13 18:54:02,186 INFO session.SessionState: Updating thread name to 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1 main
2024-09-13 18:54:02,189 INFO ql.Driver: Compiling command(queryId=monid_20240913185402_6dd7d7a1-20a7-4a58-9f3c-23359b5a2d89): SELECT id, name, age, salary FROM employees_table WHERE salary > 60000
2024-09-13 18:54:02,203 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
2024-09-13 18:54:02,203 INFO parse.CalcitePlanner: Starting Semantic Analysis
2024-09-13 18:54:02,204 INFO parse.CalcitePlanner: Completed phase 1 of Semantic Analysis
```

```
2024-09-13 18:54:02,969 INFO exec.SelectOperator: RECORDS_OUT_OPERATOR_SEL_2:4, RECORDS_OUT_INTERMEDIATE:0,
2024-09-13 18:54:02,969 INFO exec.ListSinkOperator: RECORDS_OUT_OPERATOR_LIST_SINK_5:4, RECORDS_OUT_INTERMEDIATE:0,
4      Bob Brown      35      70000.0
5      Charlie Davis  40      80000.0
7      Frank Black    32      65000.0
10     Isabel Blue    33      73000.0
Time taken: 0.76 seconds, Fetched: 4 row(s)
2024-09-13 18:54:02,978 INFO CliDriver: Time taken: 0.76 seconds, Fetched: 4 row(s)
2024-09-13 18:54:02,978 INFO conf.HiveConf: Using the default value passed in for log id: 2f3d9a62-b716-4df7-be0c-fc9f57
```

PROJECTION: (Selecting Specific Columns)

```
SELECT id, name FROM employees_table;
```

```
hive> SELECT id, name FROM employees_table;
2024-09-13 18:55:58,584 INFO conf.HiveConf: Using the default value passed in for log id: 2f3d9a62-b716-4df7-be0c-fc9f57
7d9ed1
2024-09-13 18:55:58,584 INFO session.SessionState: Updating thread name to 2f3d9a62-b716-4df7-be0c-fc9f577d9ed1 main
2024-09-13 18:55:58,586 INFO ql.Driver: Compiling command(queryId=monid_20240913185558_11f43d38-aa08-44a6-9cf6-05e8390dc
e70): SELECT id, name FROM employees_table
2024-09-13 18:55:58,598 INFO ql.Driver: Concurrency mode is disabled, not creating a lock manager
```

```
2024-09-13 18:55:58,743 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_3:10,
1      John Doe
2      Jane Smith
3      Alice Johnson
4      Bob Brown
5      Charlie Davis
6      Eve White
7      Frank Black
8      Grace Green
9      Henry Gold
10     Isabel Blue
Time taken: 0.143 seconds, Fetched: 10 row(s)
2024-09-13 18:55:58,755 INFO CliDriver: Time taken: 0.143 seconds, Fetched: 10 row(s)
```

AGGREGATION: (e.g., Summing Salaries by Age Group)

```
SELECT age, MAX(salary) AS max_salary FROM employees_table GROUP BY age;
```

22	48000.0
25	60000.0
27	52000.0
28	55000.0
29	59000.0
30	50000.0
32	
33	73000.0
35	70000.0
40	80000.0

REMOVE: (Remove Specific Records)

```
SELECT * FROM employees_table WHERE salary > 70000;
```

```
2024-09-13 19:04:58,754 INFO exec.ListSinkOperator: RECORDS_OUT_OPERATOR_LIST_SINK_5:2, RECORDS_OUT_OPERATOR_LIST_SINK_5:2
5 Charlie Davis 40 80000.0
10 Isabel Blue 33 73000.0
Time taken: 0.199 seconds, Fetched: 2 row(s)
2024-09-13 19:04:58,761 INFO cli: Time taken: 0.199 seconds, Fetched: 2 row(s)
```

COUNT: (Counting the Number of Records)

```
SELECT COUNT(*) FROM employees_table;
```

```
10
Time taken: 55.015 seconds , Fetched : 1 row(s)
```

LIMIT: (Restrict the Number of Rows Returned)

```
SELECT * FROM employees table LIMIT 5;
```

```
2024-09-13 19:09:34.014 INFO mapred.FileInputFormat: Total input files to process : 1
1      John Doe      30      50000.0
2      Jane Smith    25      60000.0
3      Alice Johnson 28      55000.0
4      Bob Brown     35      70000.0
5      Charlie Davis 40      80000.0
2024-09-13 19:09:34.031 INFO exec.TableScanOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_TS_0:5,
2024-09-13 19:09:34.032 INFO exec.SelectOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_SEL_1:5.
```

SKIP: (Skipping the First N Rows, using Row Number)

```
SELECT * FROM ( SELECT *, ROW_NUMBER() OVER () AS row_num FROM
employees table ) temp WHERE row_num > 3;
```



```

2024-09-13 18:54:02,969 INFO exec.SelectOperator: RECORDS_OUT_OPERATOR_SEL_2:4, RECORDS_OUT_INTERMEDIATE:0,
2024-09-13 18:54:02,969 INFO exec.ListSinkOperator: RECORDS_OUT_OPERATOR_LIST_SINK_5:4, RECORDS_OUT_INTERMEDIATE:0,
4      Bob Brown      35      70000.0
5      Charlie Davis   40      80000.0
7      Frank Black     32      65000.0
10     Isabel Blue     33      73000.0
Time taken: 0.76 seconds, Fetched: 4 row(s)
2024-09-13 18:54:02,978 INFO CliDriver: Time taken: 0.76 seconds, Fetched: 4 row(s)
2024-09-13 18:54:02,978 INFO conf.HiveConf: Using the default value passed in for log.id: 2f2d0a62-b716-4d57-ba9c-fc0f57

```

SORT: (Order the Data by Salary)

```
SELECT * FROM employees_table ORDER BY salary DESC;
```

```

2024-09-13 18:55:58,743 INFO exec.ListSinkOperator: RECORDS_OUT_INTERMEDIATE:0, RECORDS_OUT_OPERATOR_LIST_SINK_3:10,
1      John Doe
2      Jane Smith
3      Alice Johnson
4      Bob Brown
5      Charlie Davis
6      Eve White
7      Frank Black
8      Grace Green
9      Henry Gold
10     Isabel Blue
Time taken: 0.143 seconds, Fetched: 10 row(s)
2024-09-13 18:55:58,755 INFO CliDriver: Time taken: 0.143 seconds, Fetched: 10 row(s)

```

Result:

Thus, to import a JSON file from the command line and apply the following actions with the data present in the JSON file where, projection, aggregation, remove, count, limit, skip and sort was completed successfully.